

R Notebook

Sean Bradford

```
rm(list=ls())  
  
library(ISLR)  
library(cluster)  
library(factoextra)
```

```
## Warning: package 'factoextra' was built under R version 4.2.1
```

```
## Loading required package: ggplot2
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
library(Rfast)
```

```
## Warning: package 'Rfast' was built under R version 4.2.2
```

```
## Loading required package: Rcpp
```

```
## Warning: package 'Rcpp' was built under R version 4.2.1
```

```
## Loading required package: RcppZiggurat
```

```
## Warning: package 'RcppZiggurat' was built under R version 4.2.2
```

```
library(analogue)
```

```
## Warning: package 'analogue' was built under R version 4.2.2
```

```
## Loading required package: vegan
```

```
## Warning: package 'vegan' was built under R version 4.2.2
```

```
## Loading required package: permute
```

```
## Warning: package 'permute' was built under R version 4.2.2
```

```
## Loading required package: lattice
```

```
## This is vegan 2.6-4
```

```
## analogue version 0.17-6
```

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 4.2.1
```

```
## Registered S3 methods overwritten by 'pROC':
```

```
##   method      from
```

```
##   print.roc analogue
```

```
##   plot.roc  analogue
```

```
##
```

```
## Attaching package: 'caret'
```

```
## The following object is masked from 'package:vegan':
```

```
##
```

```
##      tolerance
```

```
cereal=read.csv("C:\\Users\\Sean\\OneDrive\\Desktop\\Grad School\\Machine Learning\\Module 8 - Hierarch
```

```
rownames(cereal)=cereal$name
```

```
cereal=cereal[,-1]
```

```
head(cereal)
```

```
##              mfr type calories protein fat sodium fiber carbo
## 100%_Bran      N  C      70      4  1    130  10.0   5.0
## 100%_Natural_Bran Q  C     120      3  5     15   2.0   8.0
## All-Bran      K  C      70      4  1    260   9.0   7.0
## All-Bran_with_Extra_Fiber K  C     50      4  0    140  14.0   8.0
## Almond_Delight R  C     110      2  2    200   1.0  14.0
## Apple_Cinnamon_Cheerios G  C     110      2  2    180   1.5  10.5
##              sugars potass vitamins shelf weight cups  rating
## 100%_Bran      6    280      25     3      1 0.33 68.40297
## 100%_Natural_Bran 8    135       0     3      1 1.00 33.98368
## All-Bran      5    320      25     3      1 0.33 59.42551
## All-Bran_with_Extra_Fiber 0    330      25     3      1 0.50 93.70491
## Almond_Delight  8     NA      25     3      1 0.75 34.38484
## Apple_Cinnamon_Cheerios 10     70      25     1      1 0.75 29.50954
```

```
# columns 1,2,12 are categorical and need to be removed before normalization
```

```
norm_cereal=scale(cereal[,c(-1,-2,-12)])
```

```
# Removing N/A values from data
```

```
norm_cereal=as.data.frame(na.omit(norm_cereal))
```

1. Apply hierarchical clustering to the data using Euclidean distance to the normalized measurements. Use Agnes to compare the clustering from single linkage, complete linkage, average linkage, and Ward. Choose the best method.

```
single=agnes(norm_cereal,method="single")
complete=agnes(norm_cereal,method="complete")
average=agnes(norm_cereal,method="average")
ward=agnes(norm_cereal,method="ward")
```

```
single$ac
```

```
## [1] 0.6091225
```

```
complete$ac
```

```
## [1] 0.8508357
```

```
average$ac
```


```
## [1] 0.7888569
```

```
ward$ac
```

```
## [1] 0.9088247
```

Ward is the best linkage method because it has the highest agglomerative coefficient.


```
d=dist(norm_cereal,method="euclidean")
d_ward=hclust(d,method="ward.D")
plot(d_ward,cex=0.6,hang=-1)
```




2. How many clusters would you choose?

Testing for best k value

```
plot(d_ward,cex=0.6)
rect.hclust(d_ward,k=4,border=1:4)
```



```
plot(d_ward,cex=0.6)
rect.hclust(d_ward,k=6,border=1:6)
```



```
k4=cutree(d_ward,k=4)
table(k4)
```

```
## k4
##  1  2  3  4
## 12 19 21 22
```

```
clustered.data=cbind.data.frame(norm_cereal,k4)
```

K = 4 appears to be the optimal value for clustering

3. Comment on the structure of the clusters and on their stability. Hint: To check stability, partition the data and see how well clusters formed based on one part apply to the other part. To do this: Cluster partition A Use the cluster centroids from A to assign each record in partition B (each record is assigned to the cluster with the closest centroid). Assess how consistent the cluster assignments are compared to the assignments based on all the data.

```
# Cluster partition A
nrow(norm_cereal)
```

```
## [1] 74
```


```
# Partitioning 80/20
74*0.8
```

```
## [1] 59.2
```

```
train=norm_cereal[1:60,]
test=norm_cereal[61:74,]
```

```
# Use cluster centroids from A to assign each record in partition B
d_train=dist(train,method="euclidean")
d_ward_train=hclust(d_train,method="ward.D")

plot(d_ward_train,cex=0.6,hang=-1)
rect.hclust(d_ward_train,k=4,border=1:4)
```



```
k4.train=cutree(d_ward_train,k=4)
table(k4.train)
```

```
## k4.train
## 1 2 3 4
## 21 12 18 9
```

```
train2=cbind.data.frame(train,k4.train)
```

```
c.1=colMeans(train2[train2$k4.train == "1",])
c.2=colMeans(train2[train2$k4.train=="2",])
c.3=colMeans(train2[train2$k4.train=="3",])
c.4=colMeans(train2[train2$k4.train=="4",])

centroid=rbind(c.1,c.2,c.3,c.4)
test.data.centroid=rowMins(distance(test,centroid[, -13]))
partition.centroid=c(train2$k4.train,test.data.centroid)
clustered.data=cbind(clustered.data,partition.centroid)
```

```
# Assess how consistent the cluster assignments are compared to the assignments based on all the data.
table(clustered.data$k4==clustered.data$partition.centroid)
```

```
##
## FALSE  TRUE
##      17    57
```

```
table(clustered.data$k4[61:74]==clustered.data$partition.centroid[61:74])
```

```
##
## FALSE  TRUE
##       2    12
```

```
(57/74)*100
```

```
## [1] 77.02703
```

```
(12/14)*100
```

```
## [1] 85.71429
```

Cluster assignments based on test data are 85.71% consistent, and the cluster assignments based on all data are 77.03% consistent.

4. The elementary public schools would like to choose a set of cereals to include in their daily cafeterias. Every day a different cereal is offered, but all cereals should support a healthy diet. For this goal, you are requested to find a cluster of “healthy cereals.” Should the data be normalized? If not, how should they be used in the cluster analysis?

```
# Calculate all centroids
```

```
ctroid1=colMeans(clustered.data[clustered.data$k4 == "1",])
ctroid2=colMeans(clustered.data[clustered.data$k4 == "2",])
ctroid3=colMeans(clustered.data[clustered.data$k4 == "3",])
ctroid4=colMeans(clustered.data[clustered.data$k4 == "4",])
ctroid.bind=rbind(ctroid1, ctroid2, ctroid3, ctroid4)
```



```
# View avg nutrient values across clusters
```


```
head(ctroid.bind)
```

```
##          calories    protein      fat      sodium      fiber      carbo
## ctroid1 -1.5080547  0.2629535 -0.75808029 -1.36294322  0.9152548 -0.4186917
## ctroid2  0.9433345  0.6074881  0.98066557 -0.04635267  0.4220701 -0.1784772
## ctroid3  0.2088503 -0.9331883 -0.01290349  0.10611786 -0.6631465 -0.6075921
## ctroid4 -0.1666359  0.1245569 -0.46452580  0.79007459 -0.1972548  0.8997331
##          sugars    potass    vitamins    weight      cups      rating k4
## ctroid1 -0.9956591  0.6639622 -0.6115433 -0.8447175 -0.3768781  1.6415416  1
## ctroid2  0.5228713  0.7739062  0.1491414  1.0099135 -0.5653466 -0.3057055  2
## ctroid3  1.0162649 -0.7283802 -0.1453172 -0.1967771  0.2328980 -0.9969602  3
## ctroid4 -0.8157229 -0.3425791  0.4650151 -0.1967771  0.4799337  0.2498970  4
##          partition.centroid
## ctroid1          1.000000
## ctroid2          1.789474
## ctroid3          3.000000
## ctroid4          2.681818
```

```
# Create heatmap to further compare cluster values
```

```
row.names(norm_cereal)=paste(k4," ",row.names(norm_cereal),sep=" ")
```

```
heatmap(as.matrix(norm_cereal),Colv=NA,hclustfun=hclust,col=rev(paste("gray",1:99,sep="")))
```



The heatmap and the table “ctroid.bind” indicate that cluster 1 would be the best choice for elementary public schools. Compared to the other clusters, cluster 1 has the highest rating & fiber, second-highest protein & potassium, and the lowest calories, fat, sodium and sugar. These factors make it the healthiest choice.

The 12 cereals shown below makeup cluster 1 and would become part of the school’s breakfast offering.

```
clustered.data[clustered.data$k4 == '1',]
```

##	calories	protein	fat	sodium
## 100%_Bran	-1.8929836	1.3286071	-0.01290349	-0.3539844
## All-Bran	-1.8929836	1.3286071	-0.01290349	1.1967306
## All-Bran_with_Extra_Fiber	-2.9194605	1.3286071	-1.00647256	-0.2346986
## Frosted_Mini-Wheats	-0.3532681	0.4151897	-1.00647256	-1.9046994
## Maypo	-0.3532681	1.3286071	-0.01290349	-1.9046994
## Puffed_Rice	-2.9194605	-1.4116451	-1.00647256	-1.9046994
## Puffed_Wheat	-2.9194605	-0.4982277	-1.00647256	-1.9046994
## Raisin_Squares	-0.8665066	-0.4982277	-1.00647256	-1.9046994
## Shredded_Wheat	-1.3797451	-0.4982277	-1.00647256	-1.9046994
## Shredded_Wheat_'n'Bran	-0.8665066	0.4151897	-1.00647256	-1.9046994
## Shredded_Wheat_spoon_size	-0.8665066	0.4151897	-1.00647256	-1.9046994
## Strawberry_Fruit_Wheats	-0.8665066	-0.4982277	-1.00647256	-1.7257708
##	fiber	carbo	sugars	potass
## 100%_Bran	3.29284661	-2.50878291	-0.234390576	2.57536849
## All-Bran	2.87327158	-1.99692385	-0.462771138	3.14346448
## All-Bran_with_Extra_Fiber	4.97114672	-1.74099432	-1.604673946	3.28548848
## Frosted_Mini-Wheats	0.35582142	-0.20541712	-0.006010015	0.01893653
## Maypo	-0.90290366	0.30644194	-0.919532261	-0.05207547
## Puffed_Rice	-0.90290366	-0.46134666	-1.604673946	-1.18826745
## Puffed_Wheat	-0.48332864	-1.22913525	-1.604673946	-0.69118346
## Raisin_Squares	-0.06375361	0.05051241	-0.234390576	0.16096053
## Shredded_Wheat	0.35582142	0.30644194	-1.604673946	-0.05207547
## Shredded_Wheat_'n'Bran	0.77539645	1.07423054	-1.604673946	0.58703252
## Shredded_Wheat_spoon_size	0.35582142	1.33016007	-1.604673946	0.30298453
## Strawberry_Fruit_Wheats	0.35582142	0.05051241	-0.462771138	-0.12308746
##	vitamins	weight	cups	rating k4
## 100%_Bran	-0.1453172	-0.1967771	-2.11003399	1.8321876 1
## All-Bran	-0.1453172	-0.1967771	-2.11003399	1.1930986 1
## All-Bran_with_Extra_Fiber	-0.1453172	-0.1967771	-1.37953029	3.6333849 1
## Frosted_Mini-Wheats	-0.1453172	-0.1967771	-0.09040611	1.1161895 1
## Maypo	-0.1453172	-0.1967771	0.76901001	0.8674423 1
## Puffed_Rice	-1.2642598	-3.5195485	0.76901001	1.2878220 1
## Puffed_Wheat	-1.2642598	-3.5195485	0.76901001	1.4479620 1
## Raisin_Squares	-0.1453172	-0.1967771	-1.37953029	0.9017710 1
## Shredded_Wheat	-1.2642598	-1.3265194	0.76901001	1.8202929 1
## Shredded_Wheat_'n'Bran	-1.2642598	-0.1967771	-0.64902659	2.2642977 1
## Shredded_Wheat_spoon_size	-1.2642598	-0.1967771	-0.64902659	2.1453309 1
## Strawberry_Fruit_Wheats	-0.1453172	-0.1967771	0.76901001	1.1887196 1
##	partition.centroid			
## 100%_Bran	1			
## All-Bran	1			
## All-Bran_with_Extra_Fiber	1			
## Frosted_Mini-Wheats	1			
## Maypo	1			

## Puffed_Rice	1
## Puffed_Wheat	1
## Raisin_Squares	1
## Shredded_Wheat	1
## Shredded_Wheat_'n'Bran	1
## Shredded_Wheat_spoon_size	1
## Strawberry_Fruit_Wheats	1